

Author: Atharva Joshi

Subject: Denoising Autoencoder

Implementation on FFHQ dataset:

Architecture:

```
class DenoisingAutoencoder(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Conv2d(1, 32, 3, stride=1, padding=1), # (32, 256, 256)
            nn.ReLU(True),
            nn.BatchNorm2d(32),
            nn.MaxPool2d(2, stride=2), # (32, 128, 128)

            nn.Conv2d(32, 64, 3, stride=1, padding=1), # (64, 128, 128)
            nn.ReLU(True),
            nn.MaxPool2d(2, stride=2), # (64, 64, 64)

            nn.Conv2d(64, 128, 3, stride=1, padding=1), # (128, 64, 64)
            nn.ReLU(True),
            nn.BatchNorm2d(128),
            nn.MaxPool2d(2, stride=2) # (128, 32, 32)

        )

        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(128, 64, 3, stride=2, padding=1,
output_padding=1), # (64, 64, 64)
            nn.ReLU(True),

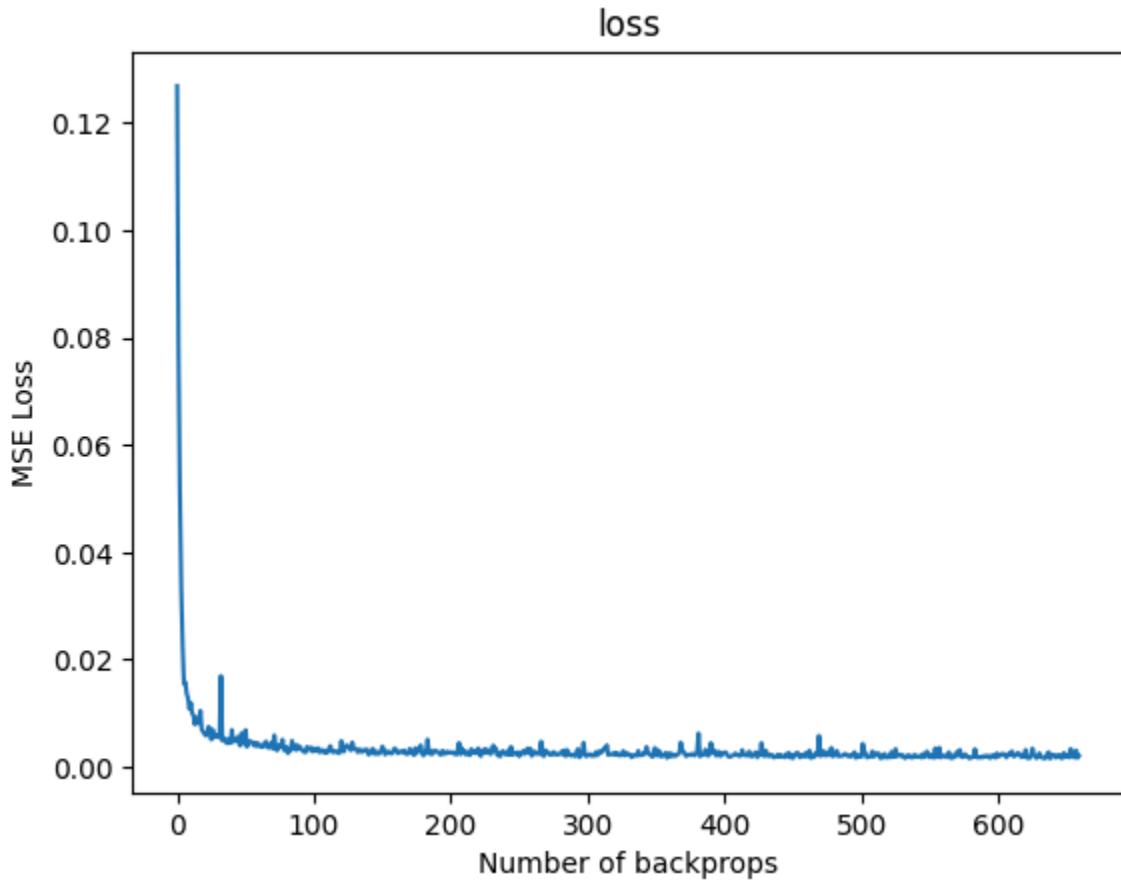
            nn.ConvTranspose2d(64, 32, 3, stride=2, padding=1,
output_padding=1), # (32, 128, 128)
            nn.ReLU(True),
            nn.BatchNorm2d(32),

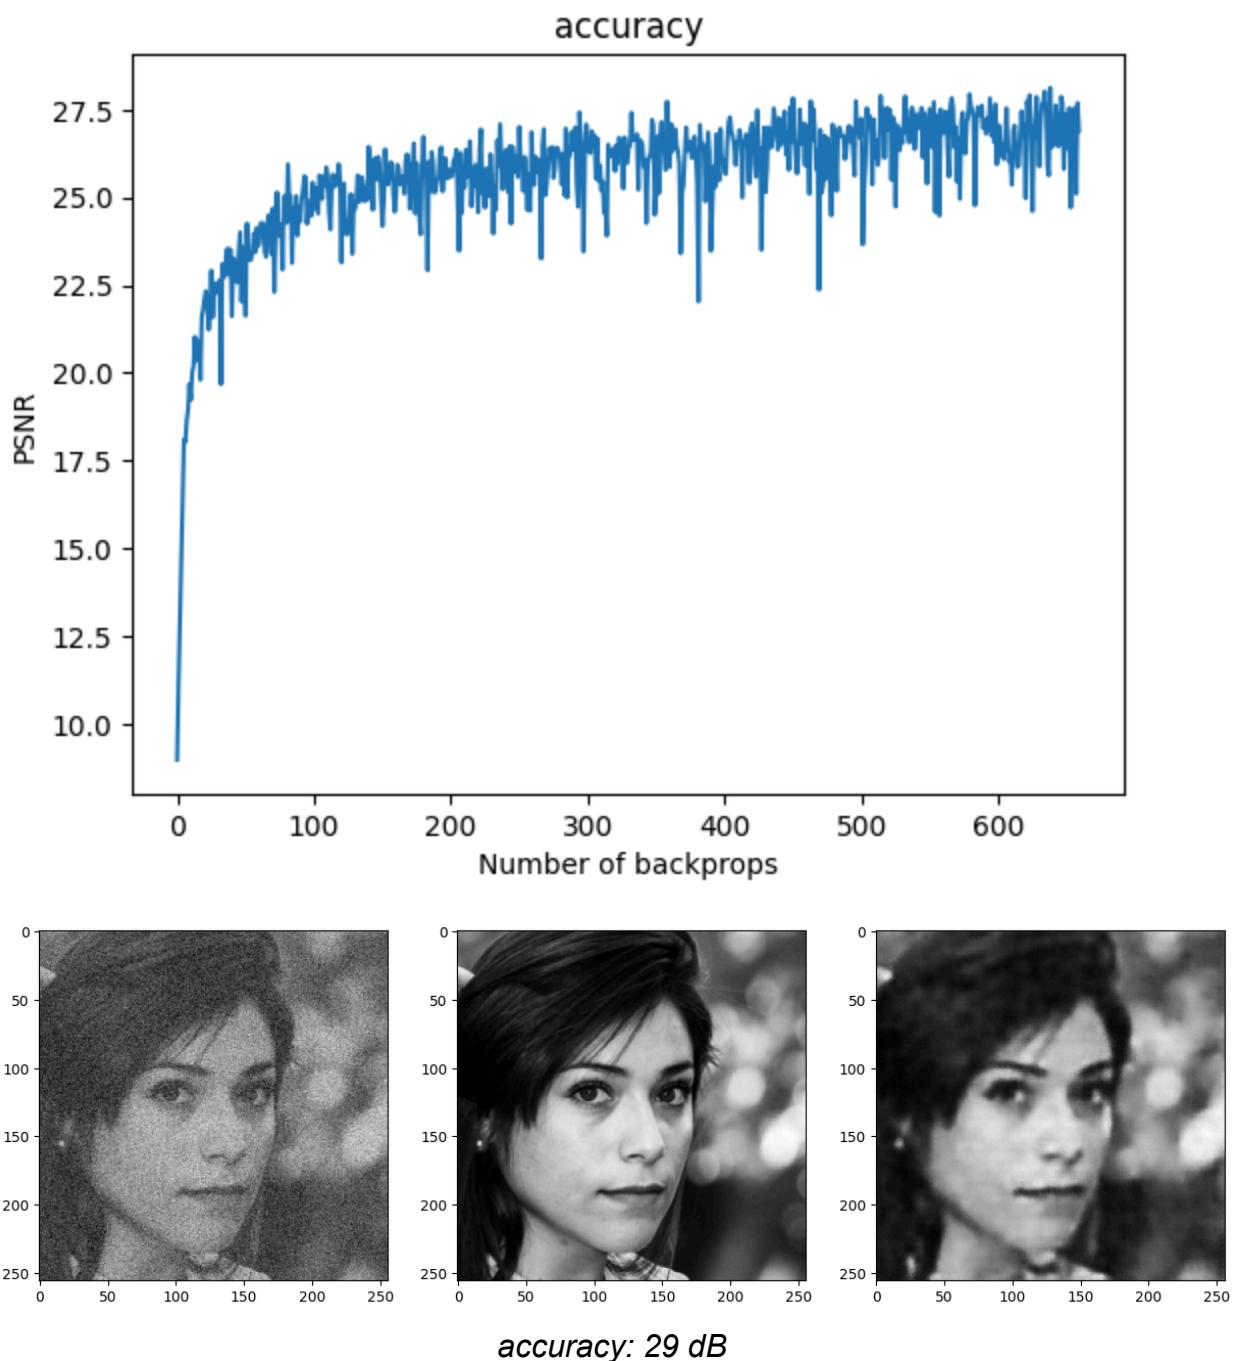
            nn.ConvTranspose2d(32, 1, 3, stride=2, padding=1,
output_padding=1), # (1, 256, 256)
            nn.Sigmoid()
        )
```

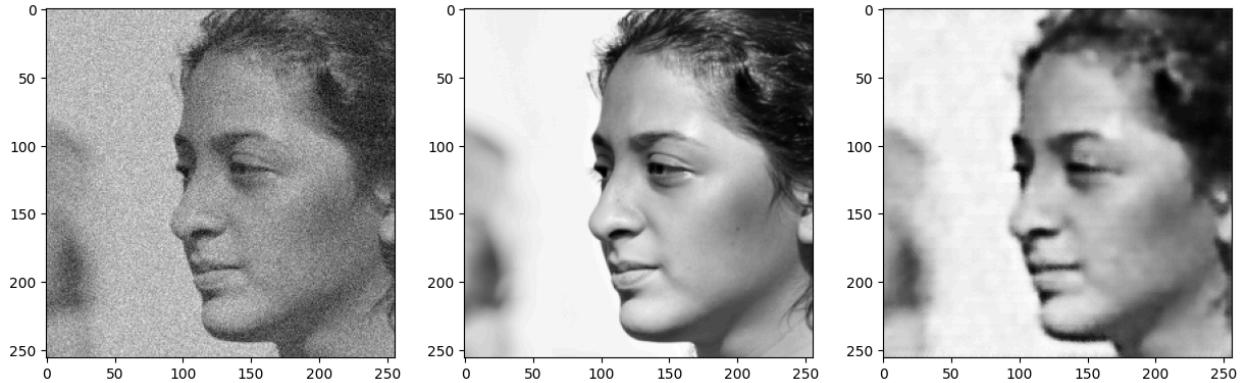
```
def forward(self, x):
    x = self.encoder(x)
    x = self.decoder(x)
    return x
```

Results:

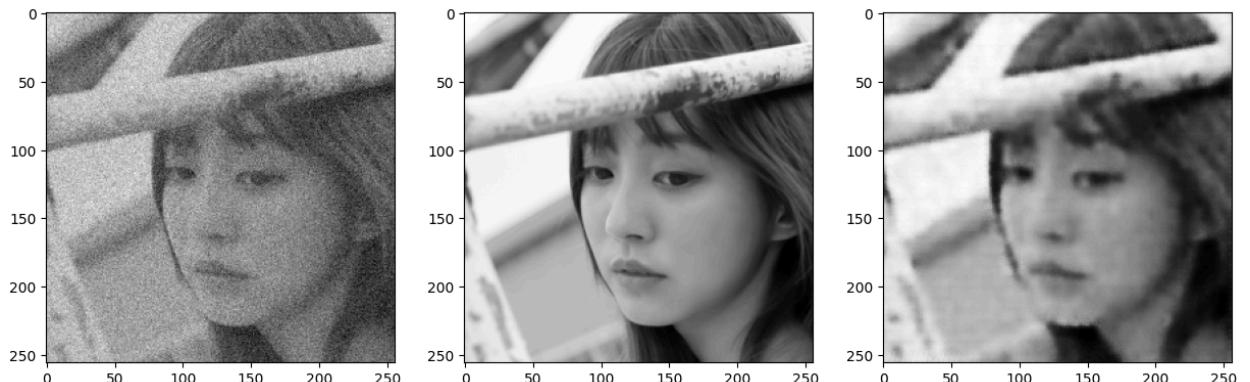
Test accuracy: 28 dB



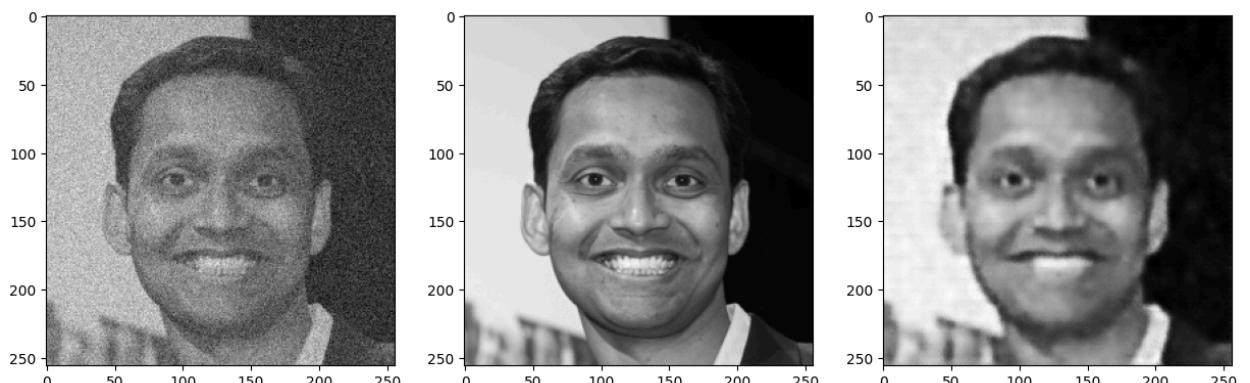




accuracy: 28 dB



accuracy: 29 dB



accuracy: 29 dB

Implementation on BSD100 dataset:

Dataset used: BSD100

Average test PSNR: 25 dB

Architecture:

```

class DenoisingAutoencoder(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Conv2d(3, 32, 3, stride=1, padding=1), # (3, 256, 256)
            nn.ReLU(True),
            nn.BatchNorm2d(32),
            nn.MaxPool2d(2, stride=2), # (32, 128, 128)

            nn.Conv2d(32, 64, 3, stride=1, padding=1), # (64, 128, 128)
            nn.ReLU(True),
            nn.MaxPool2d(2, stride=2), # (64, 64, 64)

            nn.Conv2d(64, 128, 3, stride=1, padding=1), # (128, 64, 64)
            nn.ReLU(True),
            nn.BatchNorm2d(128),
            nn.MaxPool2d(2, stride=2) # (128, 32, 32)

            )

        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(128, 64, 3, stride=2, padding=1, output_padding=1), # (64, 64, 64)
            nn.ReLU(True),

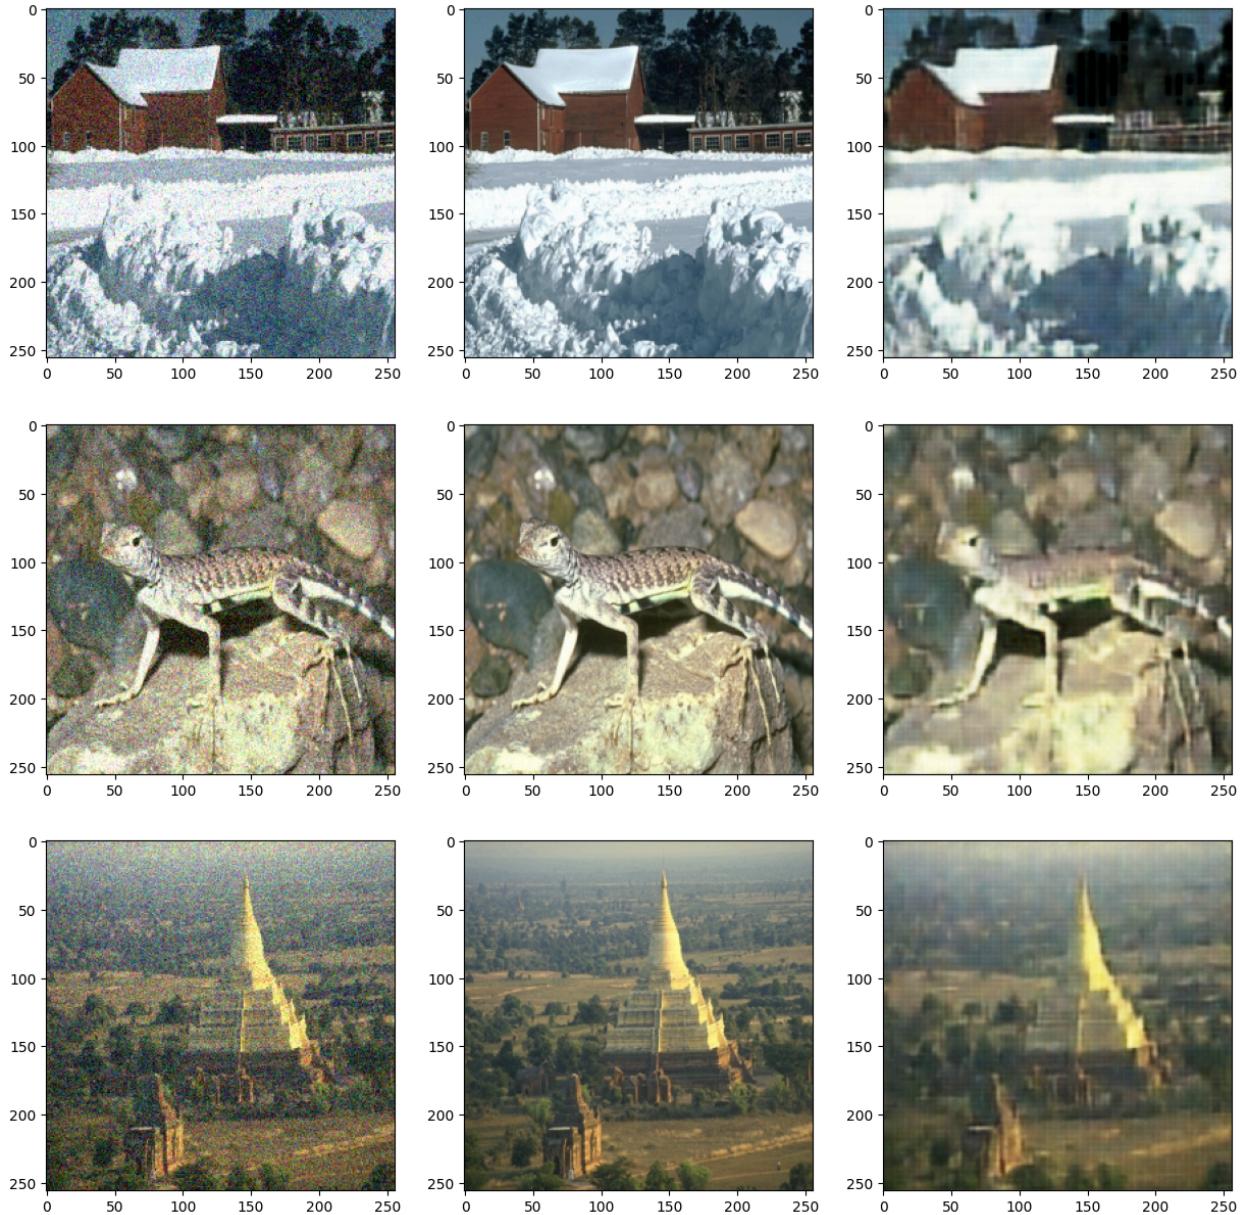
            nn.ConvTranspose2d(64, 32, 3, stride=2, padding=1, output_padding=1), # (32, 128, 128)
            nn.ReLU(True),
            nn.BatchNorm2d(32),

            nn.ConvTranspose2d(32, 3, 3, stride=2, padding=1, output_padding=1), # (3, 256, 256)
            nn.Sigmoid()
        )

    def forward(self, x):
        x = self.encoder(x)
        x = self.decoder(x)
        return x

```

Results

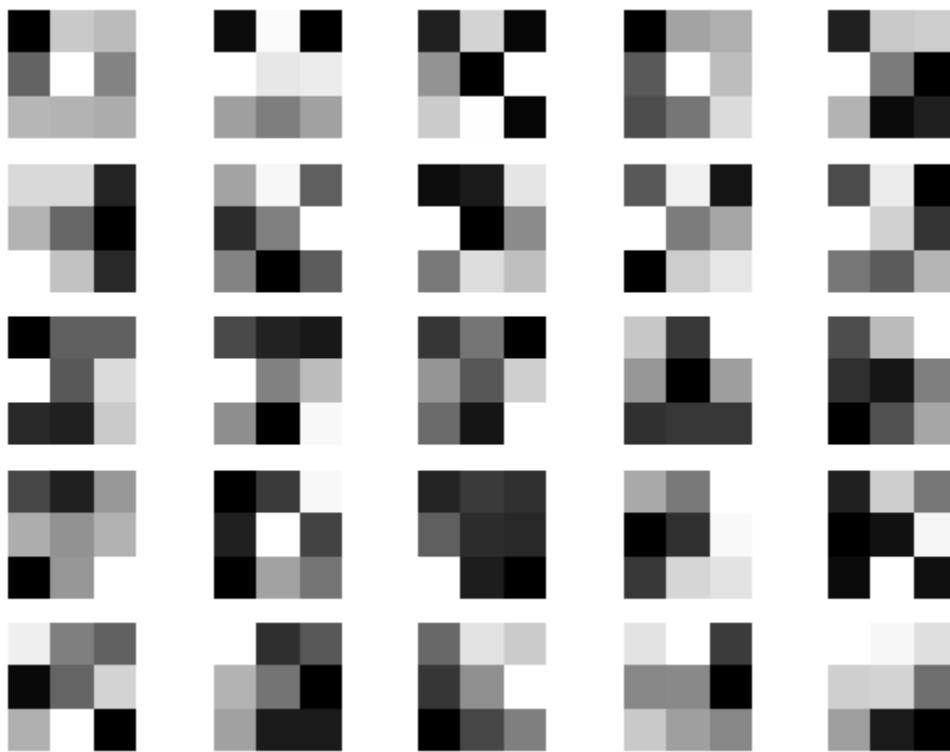


Visualising some kernels:

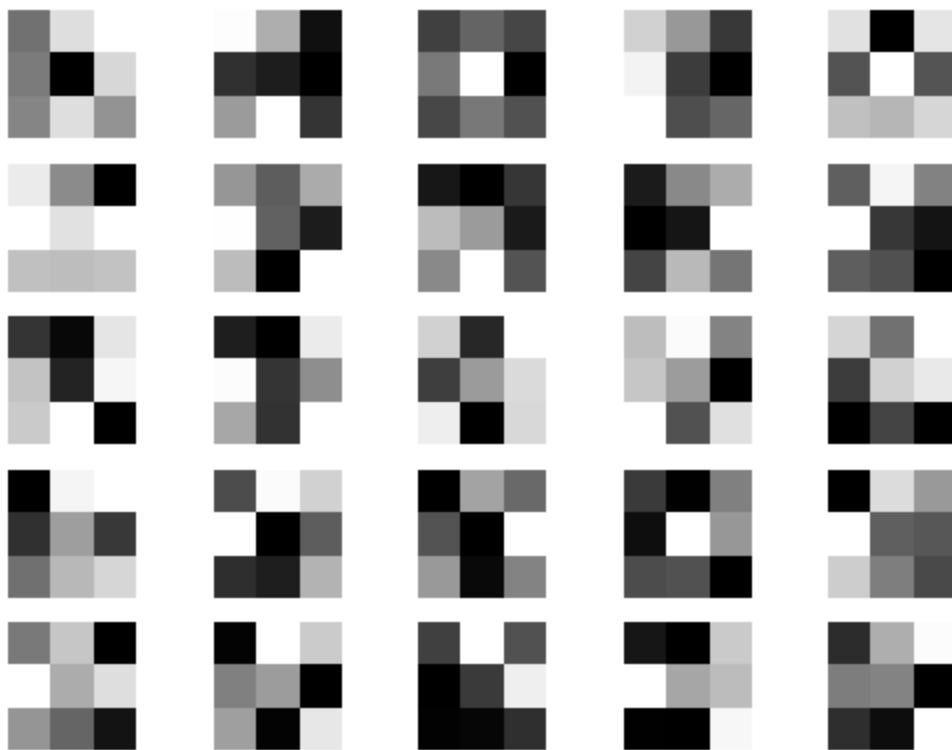
Encoder:

First layer: shape: (32, 3, 3, 3)

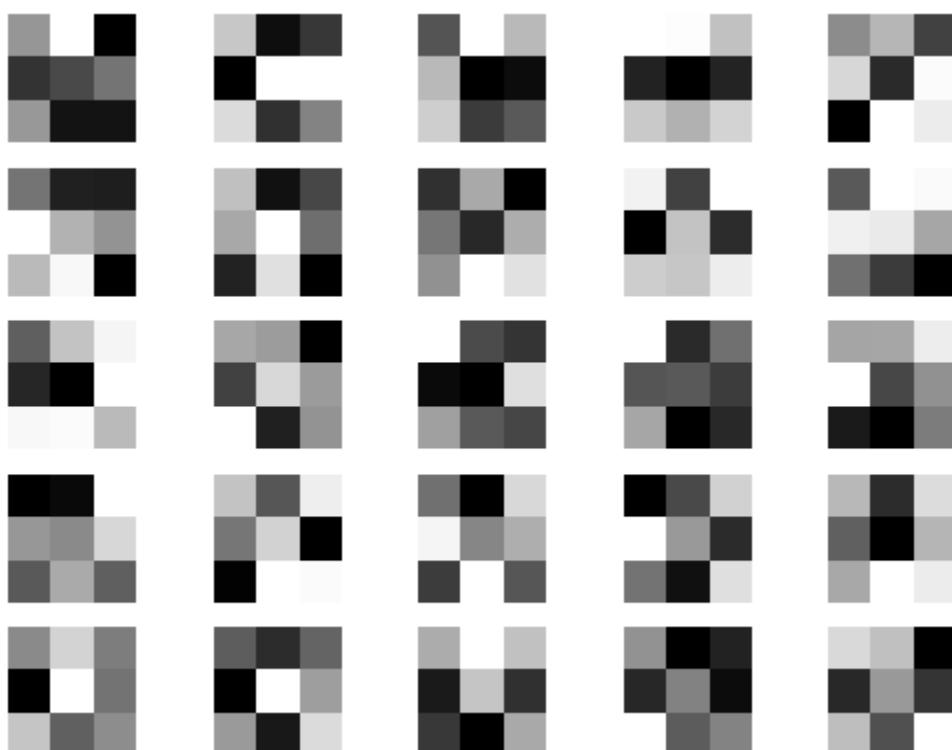
Channel 0 (Red channel):



Channel 1 (Green channel):

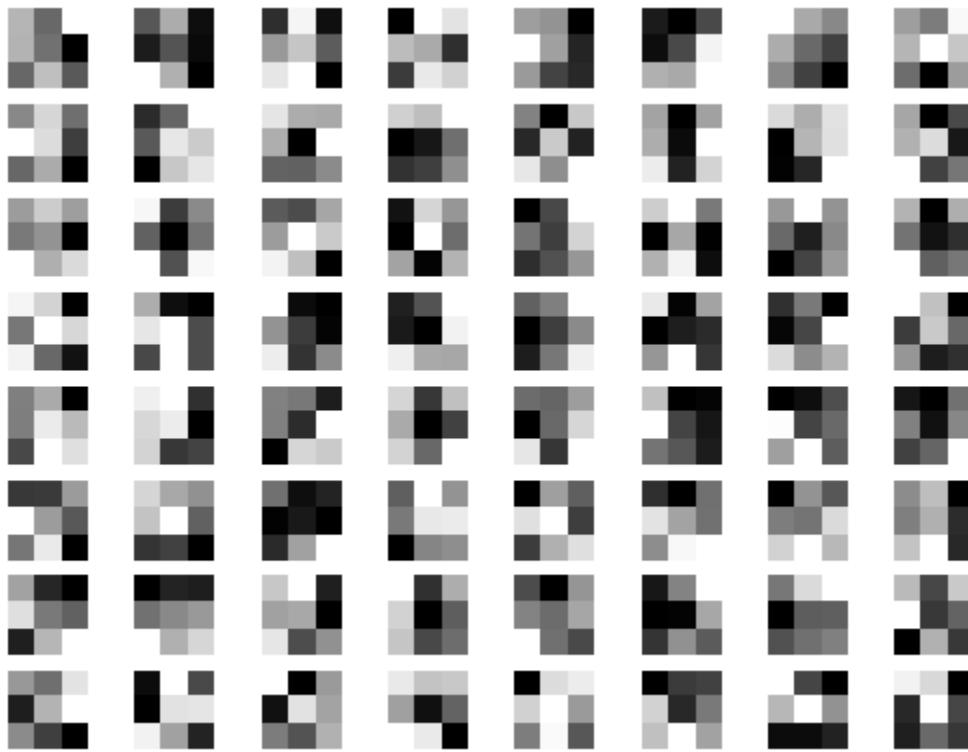


Channel 2 (Blue channel):



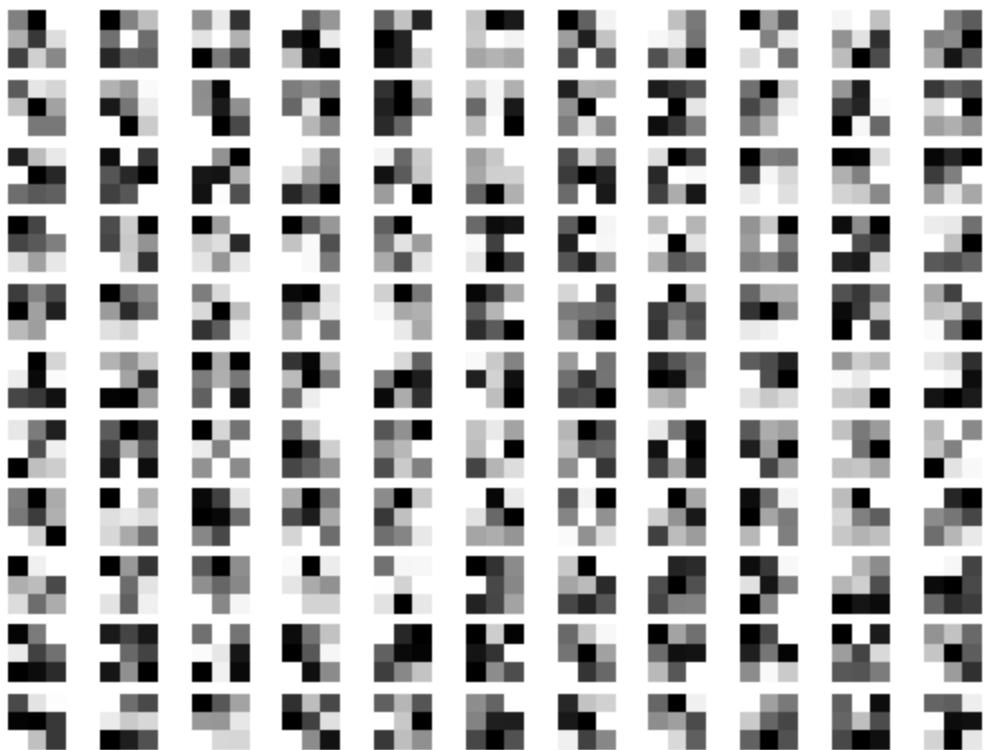
Layer 2: shape: (64, 32, 3, 3)

channel 0:



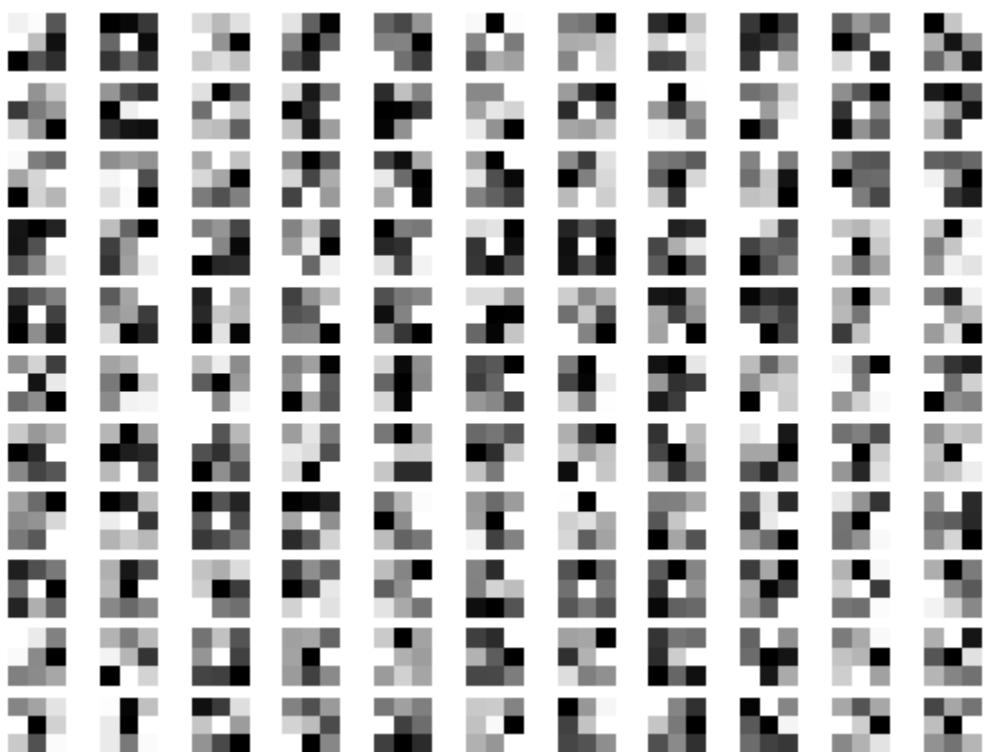
Layer 3: shape: (128, 64, 3, 3)

channel 0:

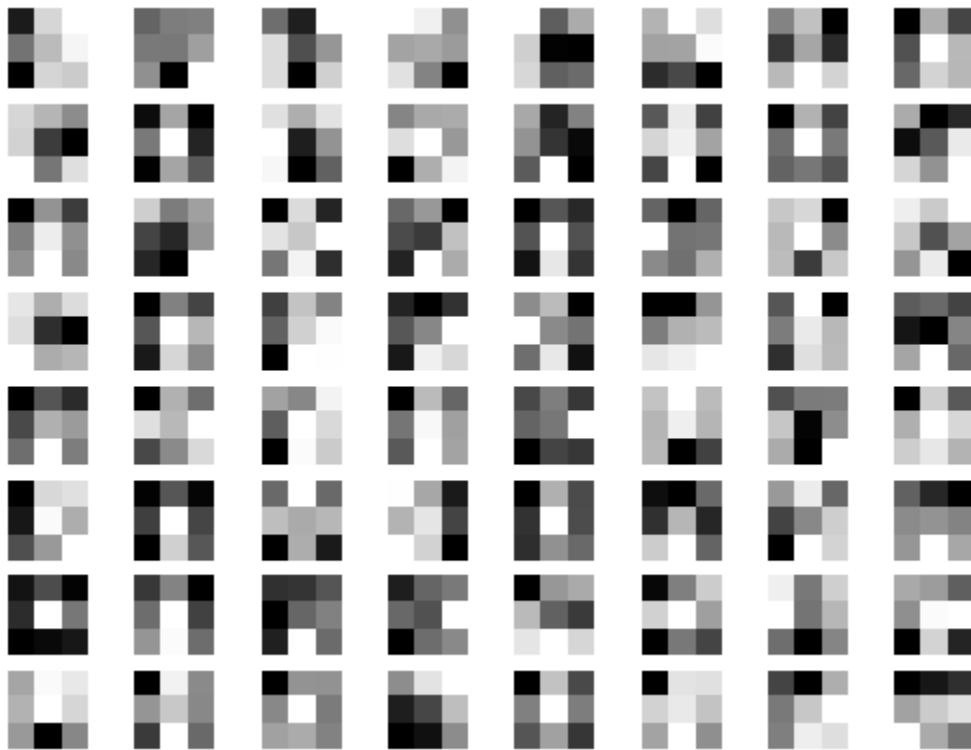


Decoder:

Layer 0:



Layer 1:



Layer 2:

